Prompt and rubric errors

HarmonyGames · Read this before you submit · Updated October 23, 2026

A task fails when someone can honestly do it two different ways, or when the rubric scores something other than the work the prompt asked for. This page shows those mistakes in plain language, with a bad line and a good line for each.

The counts come from 340 prompt reviews and 343 rubric reviews. 104 prompts failed. 130 rubrics failed. A score of 1 or 2 on any requirement fails the task. The task's score is the lowest score it got.

Date	Updates
Oct 23, 2026	Combined with the Unique Ground Truth guide: section 2 now covers what fails and what does not, the five checks a reviewer uses, and when more than one answer is acceptable. Section 6.1 has nine more patterns.
Oct 19, 2026	The clock now means your universe date: October 1, 2025 or February 14, 2026.
Oct 5, 2026	Rewritten in plain language. Same rules as the October 4 issue.
Oct 4, 2026	First issue. The checks that catch the prompt and rubric fails in those reviews, with a bad line and a good line for each.

1 · How to use this page

Go through sections 2 to 5 in order. If you cannot answer one, stop and fix the task before you continue. An empty answer means the task can still fail that check. Sections 6 and 7 are examples, for a check you are stuck on.

> **THE TEST BEFORE YOU SUBMIT**
>
> Write the outcome you want. Then write the best other outcome a careful person could still defend from the same prompt. If those two change what gets decided, what gets written, where it goes, or what it must say, find the sentence or the record that makes the other outcome impossible. If you cannot find one, the task fails.

2 · The prompt can only mean one outcome

This is the most common prompt fail: 66 of 104. Two careful people read the same prompt and do different work, and both can defend what they did. Write both outcomes here.

	The outcome you want	The other outcome someone could still defend
What gets decided		
What gets written		
Where it goes		

	The outcome you want	The other outcome someone could still defend
What it must say		
The sentence that rules the other outcome out		

“What gets written” means create, update, close, send, or leave it alone. “Where it goes” means which doc, ticket, channel, or sheet. The sentence has to be in the prompt, or in a record this person can open. “The other outcome is less likely” does not count.

These five also have to come out as one answer. If two answers fit, add a sentence to the prompt that picks one.

- Which items are included, and which are left out.
- Which number, when two amounts or two dates both exist.
- Which source wins, when two records disagree.
- What counts as the thing that happened. A slide deck, a Done ticket, or a draft is a plan, unless the prompt says that plan counts.
- What action happens. “Or” cannot mean both “close it” and “only leave a comment.”

What fails, and what does not

The score is 5 or 1–2. There is no 3 or 4. If this is the lowest score, it is the prompt score.

These do not fail. Two phrasings of the same facts. A detail the prompt never asked for. An open recommendation, when the prompt states a real boundary and every answer inside that boundary is acceptable.

These fail. Different tickets updated. A write in one reading and no write in the other. The same document with different required facts, numbers, or rows. Two homes for the same deliverable.

“Give me a recommendation” does not make every recommendation correct when the records force one answer. It also does not make one answer the only correct one when the records support several.

What a reviewer uses to pick a side

Five checks. If none of them kills the second reading, the second reading stays live.

Check	What it asks
Newer source	Is there a later record of the same fact? A later update can replace an earlier one. A missing later record does not.
Authority	Did the prompt say to follow a particular source? Quote that sentence.
Who said it	Is this the person or the file the prompt named? A different person’s words are not automatically the same fact.
Scope	Does this record settle the whole task, or only one part? A fix for one item does not decide the others.

Check	What it asks
Same result	Do both readings end with the same decision, writes, and required contents? The same destination is not enough if the contents differ.

A later rubric or oracle event cannot make an open prompt unique. Those are written after the prompt. If they require one of the two readings, they fail too.

When more than one answer is acceptable

Say so in the prompt. State the boundary in ordinary words: any recommendation grounded in the records you name, extra defects named in the same ticket, a status written in the comment or in the description.

Then the oracle events and the rubric have to accept every answer inside that boundary. They cannot quietly require one of them.

> **DO NOT FIX IT DOWNSTREAM**
>
> If the prompt leaves two live sets, a rubric row that requires one of them is wrong. Fix the prompt, or write the boundary and grade both.

3 · Five other checks on the prompt

Answer each line. If a line shows a problem, fix the prompt before you go on.

Check	What to write
Two services	Name them: _____ and _____. Each one has to change a fact the rubric scores. If you can drop one and the answer stays the same, the prompt fails.
The cause	If the prompt already says what caused the problem, delete that sentence. Keep the symptom, such as “CPI jumped,” and let the records find the cause.
Facts you stated	List every channel, ticket, vendor, amount, or date the prompt states, and the record that matches it. If the prompt states none, write “none.”
A file they cannot open	Name any record the answer needs that this person cannot open. Write “none,” or change the prompt so a missing record is an acceptable answer.
Tool names and step lists	Search for `linear_`, `gmail_`, `slack_`, `MCP`, `tool`, `First`, `Then`, and `Finally`. Delete those hits. “Check Slack” and “the Linear ticket” can stay.

Also delete any sentence you can remove without changing the task. Anything the prompt treats as already done has to be on or before your universe date: 2025-10-01 or 2026-02-14, 00:00 UTC. That is the clock inside the company data.

If you added SQL, check it. Ids must not collide with existing ones. Columns must exist. Times must fall in the author’s working hours. The author must be in the channel. An apostrophe inside a quoted string is written twice. If you added no SQL, write “no SQL.”

4 · Each result gets its own rubric row

List every result the prompt requires, including limits such as “only these rows” or “leave it open.”

For each one, write: row number, the title, a wrong answer that fails this title, another acceptable answer that still passes, and the record the fact comes from.

- No row for a result the prompt asked for: add one, on the doc, mail, ticket, or message the prompt named.
- A wrong answer still passes the title: the row is too vague. Put the real value in the title.
- Another acceptable answer fails the title: the row is too strict. Take out the channel, heading, or cause you added.
- The title says “that message,” “the above,” or “the correct figure”: paste the real value into the title.
- One row covers two writes, two services, or a verdict plus the reason: split it into two rows.
- The record is a guess: the row is wrong. Point at the message or the cell.

5 · Three searches on the rubric titles

- Search for `does not`, `must not`, `never`, `without`, `avoid`, and `fails to`. Rewrite each hit as what the agent does, or what the agent leaves in place. “The Agent reports that PR 438 had no review” can stay, because that is a fact the agent states.
- Search for `such as`, `e.g.`, and `for example`. Delete them.
- Every title starts with “The Agent” or “The Agent’s.” Fix the ones that do not. A tool name such as `linear_create_issue` does not belong in a title.

A lookup from one place is weight 1. Comparing two sources, or a small calculation, is weight 3. Combining facts from three or more sources, or applying a rule on top of a comparison, is weight 5. If one mistake fails two rows that check the same thing in the same place, delete one of them. The same figure in two different writes, such as a Slack post and a Linear ticket, is two valid rows. A repeated list keeps at most five sample rows plus one row for the count. A row must not fail every careful attempt because the value is not in the records.

One example, so the size is clear

A real prompt failed section 2. One outcome relabels only the cards that say Not Started. The other outcome relabels the whole list, because “honest again” plus the admin note also catches cards that say In Progress.

> **BAD**
> work out which of the not started ones are actually done and fix their labels the way i laid out in admin when i turned trello back on, so the board is honest again

“Not started” is one set of cards. “Honest again,” read with the admin note, can mean the whole list, including In Progress. The prompt never picks one.

> **GOOD**
> Leave In Progress and At Risk cards as they are.

That sentence makes the second outcome impossible. The rule in the admin note still has to be found.

6 · Prompt examples

Open one of these when a check in sections 2 or 3 will not come out clean. A score of 2 fails the task. Business function, competency, task skill, and the clock can stop a perfect score. They do not fail the task by themselves. Role, universe date and difficulty mode are different: a mismatch on any of them fails the task.

In plain words	How often it failed	What the fail looks like
One right outcome	66 of 104	Two different outcomes both fit the prompt
Clear about what gets written	46 of 104	The two outcomes change what gets written
Needs two services	about 40 of 104	The answer sits in one service
Stated facts are true	11 of 104	A channel, amount, or date in the prompt is wrong
This person can do it	9 of 104	This person cannot do the main request
Sounds like a colleague	8 of 104	It reads like a spec or a step list
The fact exists	4 of 104	A needed fact is not in the records
This person can open it	3 of 104	A required record is hidden from this person
The cause is left to find	1 of 104	The prompt already states the cause

These also fail the prompt: a tool name from the catalog, an extra sentence that does not change the task, a contradiction inside Snowflake or Confluence, and added SQL that breaks an id, a date, or a link.

6.1 One right outcome

This is section 2. The wording can differ. What gets decided, what gets written, where it goes, and what it must say have to be the same. A later record can pick one side. The more likely outcome does not win by itself.

> **BAD**
> work out which of the not started ones are actually done and fix their labels the way i laid out in admin when i turned trello back on, so the board is honest again

“Not started” is one set. “Honest again” plus the admin note can mean the whole list, including In Progress cards.

> **GOOD**
> On the zombie match launch list, fix the cards that still say Not Started when the work was already done. Use the rule I posted in admin the day I turned Trello back on. Leave In Progress and At Risk cards as they are.

The set is one list status. The rule still has to be found. The other cards stay out.

The patterns below come from prompt reviews that failed this check. Each one is a way two readings stay alive. The bad line is the kind of sentence that causes it. The good line is the kind of sentence that ends it. The good line still does not give away the answer.

One report fits two items

A message names a symptom, a file, or a result. More than one ticket, pull request, or card matches it.

> **BAD**
> Mark done whatever already shipped. Someone said the new faces still cannot be paired in memory or river. Go by that.

The note names no ticket. One ticket is about that symptom. A later pull request is about the memory and river screens. Attach the note to the first ticket and those screens stay done. Attach it to the pull request and they are broken. The prompt never says which match to use.

> **GOOD**
> Mark done whatever already shipped. If someone reported it broken after the merge, attach that report to the ticket whose title uses the same words. Do not attach it to a later pull request just because the files overlap.

The rule picks one ticket without naming it. A careful reader cannot put the same note on both items.

A prepared file is treated as the event

A deck, a draft, or a ticket shows that someone got ready. The prompt asks for something that happened: a conversation, a shipment, a decision.

> **BAD**
> List every publisher we talked to about this game, and where each one stands now.

Two publishers exist only as slide decks. Nothing shows an email, a message, or a reply. One reading includes them, because a deck means someone planned to talk. The other leaves them out, because “talked to” means a conversation.

> **GOOD**
> List every publisher we actually talked to about this game. A talk means an email, a Slack message, or a reply from them. A slide deck by itself does not count. Say where each one stands from the latest message.

The sentence says what counts. The deck-only names stay out. There is one list.

Two numbers, and no rule for which one

Two measurements, slices, or dates both answer the question. Each one changes the required figure or the required cut.

> **BAD**
>
> The build has to come in under 90 MB. Tell me what we still need to cut.

One report says 96 MB. A later measurement of the same build says 88 MB. Adding the same font puts one of those over the cap and the other under it. The cuts differ. The prompt does not say which measurement to use.

> **GOOD**
>
> Use the latest size measurement of the build we are actually shipping, then tell me what still has to come out to get under 90 MB.

“Latest” and “the build we are shipping” pick one row. Before you submit, check that only one measurement meets that description. If two still do, the sentence is not finished.

A scope phrase that can stretch

Words like “everything on my plate,” “the repos,” or “the open items” cover a set the records do not bound.

> **BAD**
>
> Clear everything still sitting on my plate before we shut down, and update whatever records need it.

The note that started the task is about one game. An unpaid invoice is also waiting, and it is not about that game. One reading updates only the game. The other updates the invoice too. “On my plate” covers both.

> **GOOD**
>
> Clear the items Arthur listed in his message this morning. Leave anything he did not list alone.

The set is the list in one message. The invoice stays out because it is not on that list. The prompt does not have to name the items.

The deliverable has two homes

The prompt says to write it down, track it, or post it. More than one system or channel fits.

> **BAD**
>
> Write up the open items in one tracker so we do not lose them.

Google is about to be archived. Linear is the system being kept. A tracker fits either one. One reading creates a doc. The other creates Linear issues. The prompt never picks.

> **GOOD**
> Put the open items on the existing Linear project. Do not make a Google doc. That drive is being archived.

One system is named, and the other is ruled out. The agent still has to find the project. The home is not a guess.

Two sources disagree, and nothing ranks them

An older log and a later memory, or two people, say different things about the same fact. The prompt says to report what is true and does not say which source wins.

> **BAD**
> Tell me which screenshots are actually live. Some of the old notes are stale, so go by what really happened.

The store log kept set A. A later message, written from memory, says set B is live. Nothing after the log says the live set changed, and nothing says it stayed. “What really happened” does not pick the log over the memory. The same doc can name either set.

> **GOOD**
> Tell me which screenshots are live. Use the store log. If a later message only remembers a different set, and no upload after the log says it changed, the log wins.

The prompt ranks the sources. Memory can still be mentioned as a disagreement. It cannot replace the log. A missing later record is not a ranking. If you want the log to win, say so.

The prompt and the message it points at disagree

The prompt assigns an action to one person. The thread the agent is told to read assigns it to someone else.

> **BAD**
> I am the second signer on the settlement. Read the wind-down thread and send it out for signature.

The thread says the other founder is the second signer. The prompt says this person is. Both are in the materials the agent was told to use. One reading sends it to this person. The other sends it to the founder named in the thread.

> **GOOD**
> The wind-down thread names the second signer. Use that name, even if I assumed it was me. Send it to the person the thread names.

One source is in charge. The prompt tells the agent to ignore the assumption. Read the thread before you point at it. If the prompt is supposed to win instead, say the thread is stale and name the part to ignore.

The fact needed for one answer is not visible

The prompt asks for an exact figure. The record is missing, or only another person can see it. Different substitutes then look reasonable.

> **BAD**
>
> Tell me the exact amount each founder still has to cover, including every card charge and the open invoice.

The card charges are not itemized anywhere this person can read. The invoice was emailed only to someone else. One reading guesses from a ledger total. Another leaves the invoice out. The prompt asked for one exact amount. The records do not contain it.

> **GOOD**
>
> Tell me the exact amount each founder still has to cover from the charges and invoices in my mail and the shared ledger. If a bill is missing from those, say it is missing. Do not fill it in from somewhere else.

The answer uses only records this person can open. A missing bill is part of the answer. It is not a second invented number.

Overlapping labels

The categories in the request can describe the same item in more than one way.

> **BAD**
>
> For each item, say whether it needed an approval nobody ever asked for, already has an approval nobody recorded, or has an approval that was never written down.

Those phrases overlap. The same gap can be “nobody asked” and “nobody wrote it down.” The row filed under one label is a different required row from the same gap filed under another.

> **GOOD**
>
> Use one label per item. “Not requested” means there is no ask in the thread. “Requested, not written down” means the ask is there and the approval file is empty. If both seem to fit, use “Not requested” only when no ask exists.

Each label is a check on a record. The last sentence stops one item from landing in two buckets.

6.2 Clear about what gets written

This is also section 2. Sending the same words to the same person in a different channel still passes. That is a 3, which blocks a perfect score and does not fail the task.

> **BAD**
>
> Wrap up the leftover reviews or flag them before Friday.

One outcome closes them. The other only leaves a comment.

> **GOOD**
> If a leftover review already has an owner, comment on it and leave it open. If it never shipped and nobody owns it, close it. Do not do both.

Each item has one action. The test is on the record. Use “or” only when both endings are acceptable and the later grading will accept both.

6.3 Needs two services

This is the two-services line in section 3. There is no middle score. Mentioning a second service still fails when the answer does not depend on it.

> **BAD**
> Read the #zombie-marketing thread and tell me which of those five campaigns we already paused.

The thread already lists the five and which were paused.

> **GOOD**
> Tell me which of the campaigns Jane named in #zombie-marketing are actually paused. The thread is what she remembers. The campaign record is the status.

Neither source alone is enough, and the prompt says which one wins.

6.4 Stated facts are true

This is the “facts you stated” line in section 3. One wrong channel, ticket, vendor, amount, date, or fiscal period fails the prompt. Two loose errors, such as a wrong first name, also fail it. One loose error is a 3, which blocks a perfect score.

> **BAD**
> Leonard’s note in #ua-spend on February 11 said the Adjoe test starts the 12th. Use that.

The note is in #zombie-marketing, dated February 10, and the test had already started.

> **GOOD**
> Use the note I wrote when I turned the Adjoe test on. If my date or channel is off, use the message that is actually there.

The prompt gives a lead. The record is allowed to correct it.

6.5 This person can do the task

This is the “file they cannot open” line in section 3. Gmail, Slack, calendar, Drive, Docs, Sheets, and Slides show only what this person can see. Contacts, GitHub, Trello, and Linear are open to anyone in the company. Check Confluence and Snowflake from this person’s view before you rely on them. Looking at the files yourself does not let the agent see a mailbox this person cannot open.

> **BAD**
> Pull the invoice Jane sent Marcus and reply from my mailbox with the amount.

The invoice is only in Marcus’s mailbox.

> **GOOD**
> Use the invoices in my mail and the shared ledger. If a bill is missing from those, say it is missing.

A missing bill stays missing. The prompt says so.

6.6 Tool names, step lists, and extra sentences

This is the search line in section 3. A tool name from the catalog fails. An MCP server name, or “use the Gmail tool,” is a 3, which blocks a perfect score and does not fail the task. “Check Slack” stays. A “first / then / finally” list fails. A sentence you can delete without changing the task fails. A raw id, or a time like “January 15 at 3:47 PM,” fails unless the task is about that exact timestamp.

> **BAD**
> First search Gmail for the invoice, then open Snowflake, then use linear_create_issue. Also book a dentist for Thursday.

> **GOOD**
> The Adjoe invoice and the install ledger disagree. Put the difference on the existing Linear project and tell me which number you used.

6.7 The cause is left to find

This is the cause line in section 3. There is no middle score. Naming the symptom can pass. Stating the cause fails.

> **BAD**
> The CPI jump is the Adjoe incentivized test. Write that up for Robert.

> **GOOD**
> CPI jumped last week and I don’t know which test did it. Tell Robert what the records support, and say so if they don’t settle it.

6.8 Clock, labels, broken data, and SQL

Clock. Today inside the company data is your universe date: 2025-10-01 or 2026-02-14, 00:00 UTC. Asking what already happened in May 2027 is a 3, which blocks a perfect score and does not fail the task. A future deadline can stay a deadline.

Labels on the task. Business function, competency, and task skill have no failing score. A mismatch is a 3, which blocks a perfect score. Role, universe date and difficulty mode do fail on a mismatch. Change the label if the work is a stretch. The job title does not have to match the competency.

Missing fact. A key fact that is absent, or that no tool can retrieve, fails the prompt. A detail the prompt does not need can stay missing.

Snowflake and Confluence. A closed ticket that still has later activity, or line items that do not add up to the total, fails the task even if no agent has missed it yet.

Added SQL. Skip this only when the task is saved as having no added data. An id that collides, a missing column, or a date that breaks the thread fails the task.

7 · Rubric examples

These counts are from 286 reviews that were not a perfect score. Count each row once, at its highest severity. A missing outcome counts against the rubric, and it is not added to the number of rows. If one mistake fails two rows, count that mistake once. The rubric fails when more than 10% of rows are major, more than 15% are moderate or major, or more than 20% have any issue.

Issue	Reviews	Severity
The same mistake is scored twice	128	Moderate, counted once
A required result has no row	100	Major
The row demands a detail the records do not settle	92	Major
The row states a fact the records do not contain	87	Major
Two jobs are on one row	85	Major
The row is in the wrong category	72	Moderate
The title does not contain the real value	63	Major
The row requires work the prompt did not ask for	61	Major
The row is too vague to fail a wrong answer	45	Moderate
An exact number that is the only right number	16	Minor
A grader cannot see whether the row passed	8	Major
The title is not about the agent	6	Major, and fails the wording check
An example in the title is vague	5	Moderate

These fail the rubric on their own. Weights are off by two levels on 15% of rows, or off by one level on 30%. Two rows fail every careful attempt because the value is not in the records. There are zero outcome rows. More than half of the rows are process rows. Two process rows are invalid. The same fact is graded in three or more rows. One title is not about the agent, or it names a tool. One title says what the agent must not do. A row is missing its title, category, weight, justification, or evidence.

7.1 A required result has no row

This is a missing row in section 4.

> **BAD**
> The prompt says Arthur's cut mail carries only his rows. The rubric grades the rows and never grades that the mail has nothing else.

A run that attaches the whole pack passes every row.

> **GOOD**
> The Agent sends arthur.blake@harmonygames.co a cut mail whose body contains only the eight cut rows.

Say the limit as what the body contains.

7.2 The row demands a detail the records do not settle

This is the "another acceptable answer" test in section 4. An exact value that is the only correct value is fine. `214 MB` and `214.0 MB` are the same number. Leave out a channel, a heading, or a cause the records do not settle.

> **BAD**
> The Agent's Slack post says the CPI jump was caused by the Adjoe incentivized UA test.

The records show the test was happening at the same time. They do not show it was the cause. A careful answer that refuses to name a cause fails this row.

> **GOOD**
> The Agent reports that the records do not establish a cause for the CPI jump, and names the Adjoe test as concurrent.

7.3 The row states a fact the records do not contain

This is the record test in section 4.

> **BAD**
> The Agent states that Firebase moved to the paid plan and the free tier was switched off on December 27, 2025.

Nothing in the records says that.

GOOD

The Agent states the Firebase plan status recorded on the billing page, and does not supply a switch-off date the page does not give.

7.4 Two jobs are on one row

This is section 4. Different writes and different services are separate rows. Fields inside one message can stay together. Fixing a packed row by weakening it to “at least N” or “any one of” still fails.

BAD

The Agent posts the variance to Slack and creates the Linear ticket.

GOOD

Two rows. One for the Slack post and the variance it must contain. One for the Linear ticket and the same variance.

7.5 The title does not contain the real value

This is the title test in section 4.

BAD

The Agent includes the 15 taxonomy events.

Made-up names still pass this title. A note that accepts more than the title has the same problem. A note that accepts less makes the row too strict.

GOOD

The Agent includes begin_level, left_level, and notification_permission, each marked implemented or not implemented.

7.6 The row requires work the prompt did not ask for

This is section 4.

BAD

The prompt says to mark what has not been paid. The rows require AppLovin, AdMob, and Meta to be marked paid out.

GOOD

The Agent marks each network the ledger shows as unpaid, and leaves the others unmarked.

7.7 Too vague, counted twice, or in the wrong category

These are sections 4 and 5.

> **BAD**
> The Agent emails Arthur his cut.

The mail can attach the whole pack and still pass. Another row does not save this one.

A wide row that repeats narrower rows is the same mistake, counted once. Two rows that both fail when one email misses one recipient are one mistake: delete one. Outcome is the normal category. A write labeled Process is in the wrong category.

7.8 How a title is worded

This is section 5.

> **BAD**
> The Slack update mentions the launch date.

The subject is the message, not the agent. One title like this fails the wording check for the whole rubric.

> **GOOD**
> The Agent's Slack update includes the launch date of 2026-02-12.

> **BAD**
> The Agent does not close ENG-123.

A title that says what the agent must not do fails, even when the prompt forbids that action.

> **GOOD**
> The Agent leaves ENG-123 open.

`such as`, `e.g.`, and `for example` in a title or note that gets scored are a moderate issue, once per row.

7.9 Weights, process rows, and rows that always fail

Weight	Work
1	One lookup or a literal from one source
3	One comparison, calculation, or small reconciliation
5	Facts combined from three or more sources, or a rule applied on top of a comparison

Weight the work in the row. A row does not get a 5 because it feels important. A row that packs two writes gets split, and each piece is weighted on its own.

A process row is only for an order that every valid path must follow, and that a result row cannot see. Most tasks need none. Two invalid process rows fail the rubric.

Two rows that every careful attempt will fail, because the value is not in the records or only another person can see it, fail the rubric before anyone runs the task. An agent missing a fact that is actually there is a different problem.

CONFIDENTIAL | rashmitasahoo2002@gmail.com+outlier | 2026-11-04T12:24:09.695Z